



JUNE 2026

Adaca × StratAlliance

AI Deployer Baseline Diagnostic. MoSCoW priorities and feature roadmap

HOW WE PRIORITISED

MoSCoW, in plain terms

A simple, widely used way to rank features by how essential they are. Every item is sorted into one of four buckets, so everyone agrees on what gets built first and what can wait. The name is just the four categories: Must, Should, Could and Won't.

BUILD FIRST

Must have

Essential. The product cannot launch without these.

BUILD NEXT

Should have

Important and high value, but it can still launch without them.

BUILD IF THERE'S TIME

Could have

Desirable extras, included only if time and budget allow.

NOT BUILDING NOW

Won't have

Deliberately left out for now, though they may matter later.

NEXT The next four slides take each bucket in turn.

BUILD FIRST

Must Have

The features the tool cannot launch without. If any one of these is missing, it either does not work or is not worth releasing, so they are built before anything else.

■ Form Submission

Uploads with no login needed

Staff submit documents without an account or administrator sign-in. Accepts PDF and DOCX, up to 8 files at a time, with guidance on what to include.

Automatic text reading

The tool pulls the text out of every uploaded file on its own, so nothing has to be retyped or pasted in.

Document storage is a setting

An admin decides whether uploaded files are kept, held for a set time, or deleted as soon as the assessment is done.

■ Assessment

Judged on evidence only

Each control is rated on what the documents actually show: fully covered, partly covered, or not covered.

Checked against your standard

Documents are assessed against your Responsible AI Standard, six principles and around 42 controls, all editable in the app.

Catches gaps and contradictions

It flags where documents disagree, and tests whether the governance is strong enough to actually deploy AI, not just look right on paper.

■ Report and Next Steps

One clear verdict

A single Green, Amber or Red result, with a line stating which documents it used and what was missing.

A plain-language report

Four short parts: the verdict, governance gaps, deployment readiness, and a ranked list of next steps.

Clear next steps

Three options at the end (book a review, request a deeper assessment, get help), with the report emailed to the user.

BUILD NEXT

Should Have

Important features that add real value but are not essential on day one. The tool can launch without them, and we build them as soon as the must-haves are done.

■ Deeper Coverage

Australian regulation check

A second, editable rule set that tests documents against the Privacy Act, APRA CPS 230, the APS AI Policy and National AI Centre guidance.

Traceable findings

For each control, the tool records the exact wording it relied on, so any finding can be traced back and defended.

■ Sharing the Report

Downloadable report

Let users save the report as a PDF, as well as viewing it on screen and receiving it by email.

Shareable link

A secure link to the report that can be passed to colleagues without re-running the assessment.

■ Scale and Reuse

Usage safeguards

Limits that stop the tool being overused or abused, which also keeps running costs under control.

Reusable template

Package the rule set and report format so the same tool can serve other standards or clients without rebuilding it.

BUILD IF THERE'S TIME

Could Have

Desirable extras we would include if time and budget allow. They are not essential, and the first releases are perfectly fine without them.

■ Paid Outputs

In-depth report

A more detailed, paid report that adds specific fixes for each gap and the full control-by-control breakdown.

Readiness scorecard

A standard score for how ready an organisation is to run an AI pilot, so results can be compared consistently.

■ Intake and Workflow

Use-case intake

A way for staff to submit new AI use cases and route them through review and approval, all configured in the database rather than coded.

■ Customer Accounts

Book a meeting in-app

Let users book a review directly from the report, instead of emailing separately.

Returning accounts

Optional sign-in so customers can revisit past assessments and track how their score changes over time.

NOT BUILDING NOW

Won't Have

Features we have deliberately left out of these early releases. They may matter later, but excluding them now keeps the team focused on what ships first.

■ Benchmarking

Compare against your industry

Comparing one organisation against many others in its industry. This needs data gathered from lots of customers first, which we will not have yet.

■ Partner Onboarding

Self-serve, branded sign-up

Letting other partners sign up and run their own branded version of the tool without our involvement.

■ Legal Review

Formal legal advice

Turning the report into legal advice. We keep it as guidance, and sign-off stays with a qualified reviewer.

ALREADY IN RED

The wider platform features (the register, RED analysis, dashboards, reports, audit history and admin tools) already exist in Adaca RED today, so they are available now whatever the column.

Product Roadmap

Accelerated go-live within four weeks.

TO 1 JUL · ALIGN Meeting Confirm Must scope, baseline content and ownership, and the retention default.	1 TO 15 JUL · 2 WKS Clickable prototype (Must) The full Must journey is walkable end to end: upload, assess, report, CTAs.	15 TO 29 JUL · 2 WKS Demo-ready (Must) Real engine on real uploads, live editable baseline, hardened for a customer.	POST-DEMO · SEQUENCED Fast-follow, then next Should first (regulatory layer, auditability, hardening), then Could (the monetisation paths).
---	--	--	---

AVAILABLE NOW The Adaca RED platform underpins the paid tiers in parallel: register, RED analysis, dashboards, reports, audit history and admin editor, all already shipped.